

Fraud Detection

April 3, 2024

```
[1]: #pip install imbalanced-learn
```

```
[2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
```

```
[3]: from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import RobustScaler
from imblearn.over_sampling import SMOTE
```

```
[4]: import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Conv1D, Conv2D, Flatten, \
↳BatchNormalization, Dropout
```

```
[5]: from sklearn.metrics import classification_report
```

```
[6]: df = pd.read_csv('/Users/samantarana/Downloads/onlinefraud.csv')
```

```
[7]: df.head()
```

```
[7]:
```

	step	type	amount	nameOrig	oldbalanceOrg	newbalanceOrig	\
0	1	PAYMENT	9839.64	C1231006815	170136.0	160296.36	
1	1	PAYMENT	1864.28	C1666544295	21249.0	19384.72	
2	1	TRANSFER	181.00	C1305486145	181.0	0.00	
3	1	CASH_OUT	181.00	C840083671	181.0	0.00	
4	1	PAYMENT	11668.14	C2048537720	41554.0	29885.86	

	nameDest	oldbalanceDest	newbalanceDest	isFraud	isFlaggedFraud
0	M1979787155	0.0	0.0	0	0
1	M2044282225	0.0	0.0	0	0
2	C553264065	0.0	0.0	1	0
3	C38997010	21182.0	0.0	1	0
4	M1230701703	0.0	0.0	0	0

```
[8]: df.info
```

```
[8]: <bound method DataFrame.info of
oldbalanceOrg \
0      1  PAYMENT      9839.64  C1231006815      170136.00
1      1  PAYMENT      1864.28  C1666544295      21249.00
2      1  TRANSFER       181.00  C1305486145        181.00
3      1  CASH_OUT       181.00  C840083671        181.00
4      1  PAYMENT     11668.14  C2048537720     41554.00
...
6362615  743  CASH_OUT    339682.13  C786484425    339682.13
6362616  743  TRANSFER    6311409.28  C1529008245    6311409.28
6362617  743  CASH_OUT    6311409.28  C1162922333    6311409.28
6362618  743  TRANSFER     850002.52  C1685995037     850002.52
6362619  743  CASH_OUT     850002.52  C1280323807     850002.52

newbalanceOrig  nameDest  oldbalanceDest  newbalanceDest  isFraud \
0      160296.36  M1979787155           0.00           0.00         0
1      19384.72  M2044282225           0.00           0.00         0
2           0.00  C553264065           0.00           0.00         1
3           0.00  C38997010      21182.00           0.00         1
4      29885.86  M1230701703           0.00           0.00         0
...
6362615           0.00  C776919290           0.00     339682.13         1
6362616           0.00  C1881841831           0.00           0.00         1
6362617           0.00  C1365125890     68488.84    6379898.11         1
6362618           0.00  C2080388513           0.00           0.00         1
6362619           0.00  C873221189    6510099.11    7360101.63         1

isFlaggedFraud
0           0
1           0
2           0
3           0
4           0
...
6362615           0
6362616           0
6362617           0
6362618           0
6362619           0

[6362620 rows x 11 columns]>
```

```
[9]: df.isnull().sum()
```

```
[9]: step          0
      type          0
      amount        0
      nameOrig      0
      oldbalanceOrg  0
      newbalanceOrig 0
      nameDest       0
      oldbalanceDest 0
      newbalanceDest 0
      isFraud        0
      isFlaggedFraud 0
      dtype: int64
```

```
[10]: df.corr
```

```
[10]: <bound method DataFrame.corr of
oldbalanceOrg \
0      1  PAYMENT      9839.64  C1231006815      170136.00
1      1  PAYMENT      1864.28  C1666544295      21249.00
2      1  TRANSFER      181.00  C1305486145      181.00
3      1  CASH_OUT      181.00  C840083671      181.00
4      1  PAYMENT     11668.14  C2048537720     41554.00
...
6362615  743  CASH_OUT     339682.13  C786484425     339682.13
6362616  743  TRANSFER    6311409.28  C1529008245    6311409.28
6362617  743  CASH_OUT    6311409.28  C1162922333    6311409.28
6362618  743  TRANSFER     850002.52  C1685995037     850002.52
6362619  743  CASH_OUT     850002.52  C1280323807     850002.52

newbalanceOrig  nameDest  oldbalanceDest  newbalanceDest  isFraud \
0      160296.36  M1979787155      0.00      0.00      0
1      19384.72  M2044282225      0.00      0.00      0
2           0.00  C553264065      0.00      0.00      1
3           0.00  C38997010      21182.00      0.00      1
4      29885.86  M1230701703      0.00      0.00      0
...
6362615           0.00  C776919290      0.00     339682.13      1
6362616           0.00  C1881841831      0.00      0.00      1
6362617           0.00  C1365125890     68488.84    6379898.11      1
6362618           0.00  C2080388513      0.00      0.00      1
6362619           0.00  C873221189    6510099.11    7360101.63      1

isFlaggedFraud
0      0
1      0
2      0
3      0
```

```

4          0
...
6362615    0
6362616    0
6362617    0
6362618    0
6362619    0

```

[6362620 rows x 11 columns]>

```
[11]: df.describe()
```

```

[11]:
count      step      amount  oldbalanceOrg  newbalanceOrig \
mean    2.433972e+02  1.798619e+05  8.338831e+05  8.551137e+05
std      1.423320e+02  6.038582e+05  2.888243e+06  2.924049e+06
min      1.000000e+00  0.000000e+00  0.000000e+00  0.000000e+00
25%      1.560000e+02  1.338957e+04  0.000000e+00  0.000000e+00
50%      2.390000e+02  7.487194e+04  1.420800e+04  0.000000e+00
75%      3.350000e+02  2.087215e+05  1.073152e+05  1.442584e+05
max      7.430000e+02  9.244552e+07  5.958504e+07  4.958504e+07

count      oldbalanceDest  newbalanceDest      isFraud  isFlaggedFraud
mean      1.100702e+06    1.224996e+06  1.290820e-03  2.514687e-06
std      3.399180e+06    3.674129e+06  3.590480e-02  1.585775e-03
min      0.000000e+00    0.000000e+00  0.000000e+00  0.000000e+00
25%      0.000000e+00    0.000000e+00  0.000000e+00  0.000000e+00
50%      1.327057e+05    2.146614e+05  0.000000e+00  0.000000e+00
75%      9.430367e+05    1.111909e+06  0.000000e+00  0.000000e+00
max      3.560159e+08    3.561793e+08  1.000000e+00  1.000000e+00

```

```

[12]: df.columns = ['Transaction_hours', 'Type', 'Transaction_Amt', 'Sender',
↳ 'Sender_Bal_bfr', 'Sender_Bal_aftr', 'Receiver', 'Receiver_Bal_bfr',
↳ 'Receiver_Bal_aftr', 'isFraud', 'isFlaggedFraud']

```

```

[13]: features = ['Transaction_hours', 'Transaction_Amt', 'Sender_Bal_bfr',
↳ 'Sender_Bal_aftr', 'Receiver_Bal_bfr', 'Receiver_Bal_aftr']

```

```

[14]: #Plotting all numerical features for distribution check

```


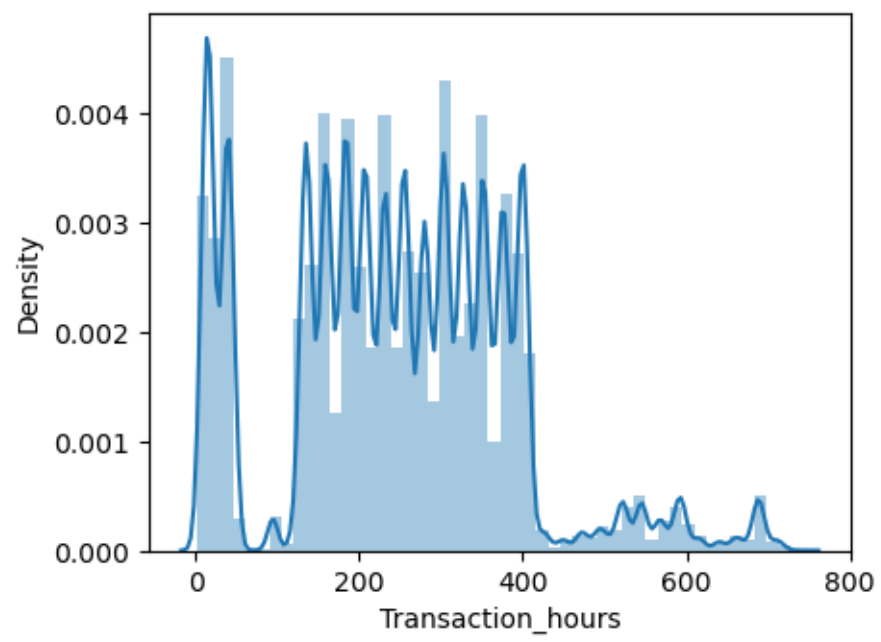
```




plt.figure(figsize=(16,8))
warnings.filterwarnings('ignore')


for feature in features:
    plt.subplot(2,3,features.index(feature)+1)
    sns.distplot(df[feature])

```

```
plt.show()
```








```
[15]: df.isFraud.value_counts(normalize=True)
```

```
[15]: 0    0.998709
      1    0.001291
      Name: isFraud, dtype: float64
```

```
[16]: df['isFraud'].value_counts().plot.pie(autopct='%1.1f%%');
```



```
[17]: non_fraud = df[df['isFraud'] == 0]
      fraud = df[df['isFraud'] == 1]
```

```
[18]: df = df.drop(['isFlaggedFraud', 'Sender', 'Receiver'], axis = 1)
```

```
[19]: df.columns
```

```
[19]: Index(['Transaction_hours', 'Type', 'Transaction_Amt', 'Sender_Bal_bfr',
        'Sender_Bal_aftr', 'Receiver_Bal_bfr', 'Receiver_Bal_aftr', 'isFraud'],
        dtype='object')
```

```
[20]: le = LabelEncoder()
      df['Type'] = le.fit_transform(df['Type'])
```

```
[21]: df.head()
```

```
[21]: Transaction_hours  Type  Transaction_Amt  Sender_Bal_bfr  Sender_Bal_aftr  \
0                    1    3          9839.64        170136.0        160296.36
1                    1    3          1864.28         21249.0         19384.72
2                    1    4           181.00           181.0             0.00
3                    1    1           181.00           181.0             0.00
4                    1    3        11668.14         41554.0        29885.86

      Receiver_Bal_bfr  Receiver_Bal_aftr  isFraud
0                0.0             0.0           0
1                0.0             0.0           0
2                0.0             0.0           1
3            21182.0             0.0           1
4                0.0             0.0           0
```

```
[22]: X = df.drop(['isFraud'], axis = 1)
```

```
[23]: Y = df.isFraud
```

```
[24]: X.shape, Y.shape
```

```
[24]: ((6362620, 7), (6362620,))
```

```
[25]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.2,
        random_state = 99)
```

```
[26]: X_train.shape, X_test.shape, Y_train.shape, Y_test.shape
```

```
[26]: ((5090096, 7), (1272524, 7), (5090096,), (1272524,))
```



```
[27]: scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
[28]: Y_train = Y_train.to_numpy()
Y_test = Y_test.to_numpy()
```

```
[29]: X_train = X_train.reshape(X_train.shape[0], X_train.shape[1])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1])
```

```
[30]: print("Before OverSampling, counts of label '1' : {}".format(sum(Y_train==1)))
print("Before OverSampling, counts of label '0' : {}".format(sum(Y_train==0)))
```

```
Before OverSampling, counts of label '1' : 6545
Before OverSampling, counts of label '0' : 5083551
```

```
[31]: #SMOTE

smote = SMOTE(random_state = 42)
X_res, Y_res = smote.fit_resample(X_train, Y_train.ravel())

X_res = X_res.reshape(X_res.shape[0], X_res.shape[1], 1)
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1], 1)
```

```
[32]: X_res.shape
```

```
[32]: (10167102, 7, 1)
```

```
[33]: print("After OverSampling, the shape of train_X: {}".format(X_res.shape))
print("After OverSampling, the shape of train_Y: {}".format(Y_res.shape))
```

```
After OverSampling, the shape of train_X: (10167102, 7, 1)
After OverSampling, the shape of train_Y: (10167102,)
```

```
[34]: print("After OverSampling, counts of label '1' : {}".format(sum(Y_res==1)))
print("After OverSampling, counts of label '0' : {}".format(sum(Y_res==0)))
```

```
After OverSampling, counts of label '1' : 5083551
After OverSampling, counts of label '0' : 5083551
```

```
[35]: model = Sequential()

#layers
model.add(Conv1D(filters=32, kernel_size = 2, activation='relu', input_shape = X_res[0].shape))
model.add(BatchNormalization())
model.add(Dropout(0.2))

model.add(Conv1D(filters=64, kernel_size=2, activation='relu'))
```

```
model.add(BatchNormalization())
model.add(Dropout(0.2))
```

```
[36]: #build ANN
model.add(Flatten())
model.add(Dense(64, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(1, activation='relu'))
```

```
[37]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv1d (Conv1D)	(None, 6, 32)	96
batch_normalization (Batch Normalization)	(None, 6, 32)	128
dropout (Dropout)	(None, 6, 32)	0
conv1d_1 (Conv1D)	(None, 5, 64)	4160
batch_normalization_1 (Batch Normalization)	(None, 5, 64)	256
dropout_1 (Dropout)	(None, 5, 64)	0
flatten (Flatten)	(None, 320)	0
dense (Dense)	(None, 64)	20544
dropout_2 (Dropout)	(None, 64)	0
dense_1 (Dense)	(None, 1)	65

```
====
Total params: 25249 (98.63 KB)
Trainable params: 25057 (97.88 KB)
Non-trainable params: 192 (768.00 Byte)
=====
```

```
[38]: model.compile(optimizer='adam', loss='binary_crossentropy',
    ↪ metrics=['accuracy'])
```

```
[39]: %%time

#fitting the model

history1 = model.fit(X_res, Y_res, epochs=1, validation_data=(X_test, Y_test))

317722/317722 [=====] - 499s 2ms/step - loss: 0.1921 -
accuracy: 0.9570 - val_loss: 0.0862 - val_accuracy: 0.9699
CPU times: user 7min 20s, sys: 3min 3s, total: 10min 24s
Wall time: 8min 19s
```

```
[40]: #import model
model12 = Sequential()

#layers
model12.add(Conv1D(filters=32, kernel_size = 2, activation='relu', input_shape_
↪ X_res[0].shape))
model12.add(BatchNormalization())
model12.add(Dropout(0.2))

model12.add(Conv1D(filters=64, kernel_size=2, activation='relu'))
model12.add(BatchNormalization())
model12.add(Dropout(0.2))
```

```
[41]: #build ANN
model12.add(Flatten())
model12.add(Dense(64, activation = 'relu'))
model12.add(Dropout(0.5))
model12.add(Dense(1, activation = 'relu'))
```

```
[42]: # convert model weights to bfloat16 data type

policy = tf.keras.mixed_precision.Policy('mixed_bfloat16')
tf.keras.mixed_precision.set_global_policy(policy)
```

```
[43]: X_train_bfloat = tf.convert_to_tensor(X_res, dtype = tf.bfloat16)
```

```
[44]: X_train_bfloat.shape
```

```
[44]: TensorShape([10167102, 7, 1])
```

```
[45]: X_test_bfloat = tf.convert_to_tensor(X_test, dtype = tf.bfloat16)
```

```
[46]: model12.compile(optimizer = 'adam', loss = 'binary_crossentropy', metrics =_
↪ ['accuracy'])
```

```
[47]: %%time
```

```
#fitting the model
```

```
history2 = model12.fit(X_train_bfloat, Y_res, epochs=1, validation_data =  
↳(X_test_bfloat, Y_test))
```

```
317722/317722 [=====] - 2296s 7ms/step - loss: 0.1810 -  
accuracy: 0.9608 - val_loss: 0.1313 - val_accuracy: 0.9275  
CPU times: user 7min 28s, sys: 3min 5s, total: 10min 33s  
Wall time: 38min 16s
```

0.0.1 With 92% of accuracy rate , we have established a working model for the fraud detection !