

Python Portfolio Project - Automatidata

February 27, 2024

0.1 *Automatidata Project*

0.1.1 Project to rationalise the pick and drop-off services of the Automatidata corporation.

```
[1]: # Importing required libraries

import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import datetime as dt
import seaborn as sns
```

```
[2]: #importing the dataset collected from the company

df=pd.read_csv('/Users/samantarana/Downloads/2017_Yellow_Taxi_Trip_Data.csv')
```

Heading ahead with 'Data Exploration' and Cleaning

```
[3]: df.head(10)
```

```
[3]:   Unnamed: 0  VendorID  tpep_pickup_datetime  tpep_dropoff_datetime  \
0    24870114         2  03/25/2017 8:55:43 AM  03/25/2017 9:09:47 AM  \
1    35634249         1  04/11/2017 2:53:28 PM  04/11/2017 3:19:58 PM  \
2    106203690        1  12/15/2017 7:26:56 AM  12/15/2017 7:34:08 AM  \
3    38942136         2  05/07/2017 1:17:59 PM  05/07/2017 1:48:14 PM  \
4    30841670         2  04/15/2017 11:32:20 PM  04/15/2017 11:49:03 PM  \
5    23345809         2  03/25/2017 8:34:11 PM  03/25/2017 8:42:11 PM  \
6    37660487         2  05/03/2017 7:04:09 PM  05/03/2017 8:03:47 PM  \
7    69059411         2  08/15/2017 5:41:06 PM  08/15/2017 6:03:05 PM  \
8     8433159         2  02/04/2017 4:17:07 PM  02/04/2017 4:29:14 PM  \
9    95294817         1  11/10/2017 3:20:29 PM  11/10/2017 3:40:55 PM  \

   passenger_count  trip_distance  RatecodeID  store_and_fwd_flag  \
0                6           3.34          1                N
1                1           1.80          1                N
2                1           1.00          1                N
3                1           3.70          1                N
4                1           4.37          1                N
```

5	6	2.30	1	N
6	1	12.83	1	N
7	1	2.98	1	N
8	1	1.20	1	N
9	1	1.60	1	N

	PULocationID	DOLocationID	payment_type	fare_amount	extra	mta_tax	\
0	100	231	1	13.0	0.0	0.5	
1	186	43	1	16.0	0.0	0.5	
2	262	236	1	6.5	0.0	0.5	
3	188	97	1	20.5	0.0	0.5	
4	4	112	2	16.5	0.5	0.5	
5	161	236	1	9.0	0.5	0.5	
6	79	241	1	47.5	1.0	0.5	
7	237	114	1	16.0	1.0	0.5	
8	234	249	2	9.0	0.0	0.5	
9	239	237	1	13.0	0.0	0.5	

	tip_amount	tolls_amount	improvement_surcharge	total_amount
0	2.76	0.0	0.3	16.56
1	4.00	0.0	0.3	20.80
2	1.45	0.0	0.3	8.75
3	6.39	0.0	0.3	27.69
4	0.00	0.0	0.3	17.80
5	2.06	0.0	0.3	12.36
6	9.86	0.0	0.3	59.16
7	1.78	0.0	0.3	19.58
8	0.00	0.0	0.3	9.80
9	2.75	0.0	0.3	16.55

```
[4]: df.size
```

```
[4]: 408582
```

```
[5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 22699 entries, 0 to 22698
Data columns (total 18 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Unnamed: 0                            22699 non-null  int64
1   VendorID                              22699 non-null  int64
2   tpep_pickup_datetime                  22699 non-null  object
3   tpep_dropoff_datetime                  22699 non-null  object
4   passenger_count                        22699 non-null  int64
5   trip_distance                          22699 non-null  float64
6   RatecodeID                            22699 non-null  int64
```

```

7   store_and_fwd_flag      22699 non-null object
8   PULocationID            22699 non-null int64
9   DOLocationID            22699 non-null int64
10  payment_type             22699 non-null int64
11  fare_amount              22699 non-null float64
12  extra                    22699 non-null float64
13  mta_tax                  22699 non-null float64
14  tip_amount               22699 non-null float64
15  tolls_amount             22699 non-null float64
16  improvement_surcharge    22699 non-null float64
17  total_amount             22699 non-null float64

```

dtypes: float64(8), int64(7), object(3)

memory usage: 3.1+ MB

```
[6]: df.describe()
```

```

[6]:      Unnamed: 0      VendorID  passenger_count  trip_distance \
count  2.269900e+04  22699.000000      22699.000000      22699.000000
mean    5.675849e+07      1.556236          1.642319          2.913313
std     3.274493e+07      0.496838          1.285231          3.653171
min     1.212700e+04      1.000000          0.000000          0.000000
25%     2.852056e+07      1.000000          1.000000          0.990000
50%     5.673150e+07      2.000000          1.000000          1.610000
75%     8.537452e+07      2.000000          2.000000          3.060000
max     1.134863e+08      2.000000          6.000000         33.960000

      RatecodeID  PULocationID  DOLocationID  payment_type  fare_amount \
count  22699.000000  22699.000000  22699.000000  22699.000000  22699.000000
mean     1.043394    162.412353    161.527997     1.336887     13.026629
std     0.708391     66.633373     70.139691     0.496211     13.243791
min     1.000000     1.000000     1.000000     1.000000    -120.000000
25%     1.000000    114.000000    112.000000     1.000000      6.500000
50%     1.000000    162.000000    162.000000     1.000000      9.500000
75%     1.000000    233.000000    233.000000     2.000000     14.500000
max     99.000000    265.000000    265.000000     4.000000    999.990000

      extra      mta_tax      tip_amount  tolls_amount \
count  22699.000000  22699.000000  22699.000000  22699.000000
mean     0.333275     0.497445     1.835781     0.312542
std     0.463097     0.039465     2.800626     1.399212
min    -1.000000    -0.500000     0.000000     0.000000
25%     0.000000     0.500000     0.000000     0.000000
50%     0.000000     0.500000     1.350000     0.000000
75%     0.500000     0.500000     2.450000     0.000000
max      4.500000     0.500000    200.000000    19.100000

      improvement_surcharge  total_amount

```


count	22699.000000	22699.000000
mean	0.299551	16.310502
std	0.015673	16.097295
min	-0.300000	-120.300000
25%	0.300000	8.750000
50%	0.300000	11.800000
75%	0.300000	17.800000
max	0.300000	1200.290000

The function info reveals that there in no missing values


0.1.2 Checking for outliers with the help of Boxplots

```
[7]: # Convert data columns to datetime
df['tpep_pickup_datetime']=pd.to_datetime(df['tpep_pickup_datetime'])
df['tpep_dropoff_datetime']=pd.to_datetime(df['tpep_dropoff_datetime'])
```


```
[8]: # Create box plot of trip_distance
plt.figure(figsize=(7,2))
plt.title('trip_distance')
sns.boxplot(data=None, x=df['trip_distance'], fliersize=1);
```




```
[9]: # Create histogram of trip_distance
plt.figure(figsize=(10,5))
sns.histplot(df['trip_distance'], bins=range(0,26,1))
plt.title('Trip distance histogram');
```



```
[10]: # Create box plot of total_amount
plt.figure(figsize=(7,2))
plt.title('total_amount')
sns.boxplot(x=df['total_amount'], fliersize=1);
```




```
[11]: # Create histogram of total_amount
plt.figure(figsize=(12,6))
ax = sns.histplot(df['total_amount'], bins=range(-10,101,5))
ax.set_xticks(range(-10,101,5))
ax.set_xticklabels(range(-10,101,5))
plt.title('Total amount histogram');
```




observation : The total cost of each trip also has a distribution that skews right, with most costs falling in the \$5-15 range

```
[12]: # Create box plot of tip_amount
plt.figure(figsize=(7,2))
plt.title('tip_amount')
sns.boxplot(x=df['tip_amount'], fliersize=1);
```




```
[13]: # Create histogram of tip_amount
plt.figure(figsize=(12,6))
ax = sns.histplot(df['tip_amount'], bins=range(0,21,1))
ax.set_xticks(range(0,21,2))
```

```
ax.set_xticklabels(range(0,21,2))
plt.title('Tip amount histogram');
```



The distribution for tip amount is right-skewed, with nearly all the tips in the \$0-3 range.

```
[14]: # Create histogram of tip_amount by vendor
plt.figure(figsize=(12,7))
ax = sns.histplot(data=df, x='tip_amount', bins=range(0,21,1),
                  hue='VendorID',
                  multiple='stack',
                  palette='pastel')
ax.set_xticks(range(0,21,1))
ax.set_xticklabels(range(0,21,1))
plt.title('Tip amount by vendor histogram');
```

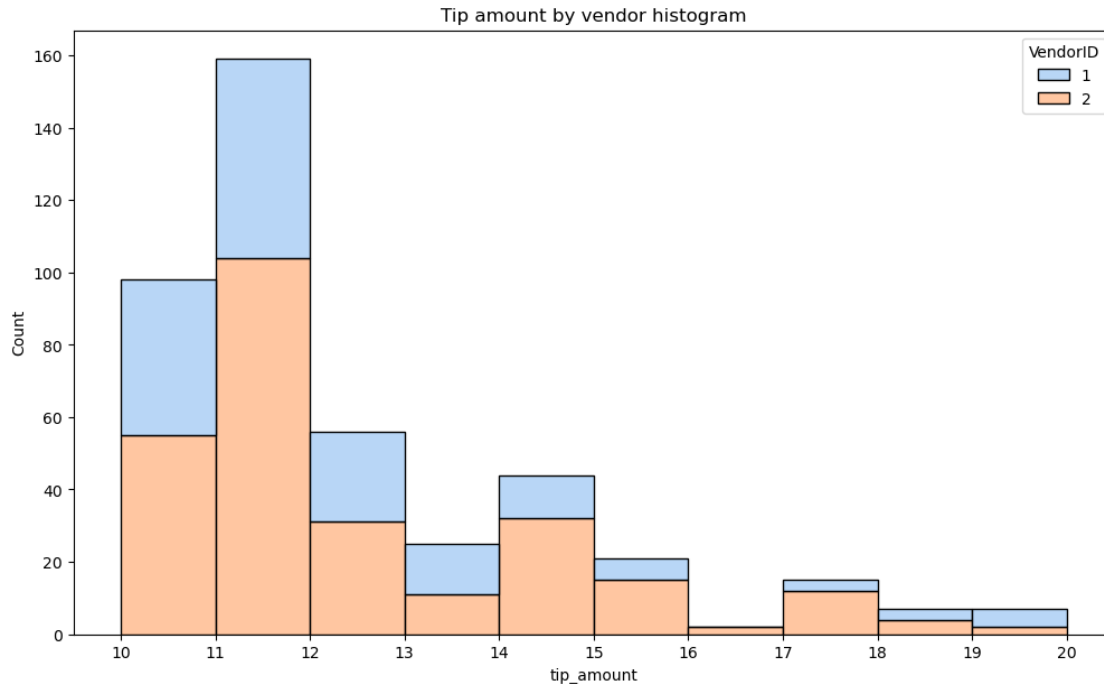


observation : Separating the tip amount by vendor reveals that there are no noticeable aberrations in the distribution of tips between the two vendors in the dataset. Vendor two has a slightly higher share of the rides, and this proportion is approximately maintained for all tip amounts.

Next, zoom in on the upper end of the range of tips to check whether vendor one gets noticeably more of the most generous tips.

```
[15]: # Create histogram of tip_amount by vendor for tips > $10

tips_over_ten = df[df['tip_amount'] > 10]
plt.figure(figsize=(12,7))
ax = sns.histplot(data=tips_over_ten, x='tip_amount', bins=range(10,21,1),
                  hue='VendorID',
                  multiple='stack',
                  palette='pastel')
ax.set_xticks(range(10,21,1))
ax.set_xticklabels(range(10,21,1))
plt.title('Tip amount by vendor histogram');
```

The proportions are maintained even at these higher tip amounts, with the exception being at highest extremity, but this is not noteworthy due to the low sample size at these tip amounts.

```
[16]: df['passenger_count'].value_counts()
```

```
[16]: 1    16117
      2     3305
      5     1143
      3      953
      6      693
      4      455
      0        33
      Name: passenger_count, dtype: int64
```


observation : Nearly two thirds of the rides were single occupancy, though there were still nearly 700 rides with as many as six passengers. Also, there are 33 rides with an occupancy count of zero, which doesn't make sense. These would likely be dropped unless a reasonable explanation can be found for them.

```
[17]: # Calculate mean tips by passenger_count
mean_tips_by_passenger_count = df.groupby(['passenger_count']).
    .mean()[['tip_amount']]
mean_tips_by_passenger_count
```

```
[17]:
```

passenger_count	tip_amount
0	2.135758
1	1.848920
2	1.856378
3	1.716768
4	1.530264
5	1.873185
6	1.720260

```
[18]: # Create bar plot for mean tips by passenger count
data = mean_tips_by_passenger_count.tail(-1)
pal = sns.color_palette("Greens_d", len(data))
rank = data['tip_amount'].argsort().argsort()
plt.figure(figsize=(12,7))
ax = sns.barplot(x=data.index,
                 y=data['tip_amount'],
                 palette=np.array(pal[::-1])[rank])
ax.axhline(df['tip_amount'].mean(), ls='--', color='red', label='global mean')
ax.legend()
plt.title('Mean tip amount by passenger count', fontsize=16);
```



Mean tip amount varies very little by passenger count. Although it does drop noticeably for four-passenger rides, it's expected that there would be a higher degree of fluctuation because rides with four passengers were the least plentiful in the dataset

(aside from rides with zero passengers).

```
[19]: # Creating a month column
df['month'] = df['tpep_pickup_datetime'].dt.month_name()

# Create a day column
df['day'] = df['tpep_pickup_datetime'].dt.day_name()
```

```
[20]: # Getting total number of rides for each month

monthly_rides = df['month'].value_counts()
monthly_rides
```

```
[20]: March          2049
      October       2027
      April         2019
      May           2013
      January       1997
      June          1964
      December      1863
      November      1843
      February      1769
      September     1734
      August        1724
      July          1697
      Name: month, dtype: int64
```

```
[21]: # Reorder the monthly ride list so months go in order
month_order = ['January', 'February', 'March', 'April', 'May', 'June', 'July',
               'August', 'September', 'October', 'November', 'December']

monthly_rides = monthly_rides.reindex(index=month_order)
monthly_rides
```

```
[21]: January          1997
      February        1769
      March           2049
      April           2019
      May             2013
      June            1964
      July            1697
      August          1724
      September       1734
      October         2027
      November        1843
      December        1863
      Name: month, dtype: int64
```


```
[22]: # Show the index
```

```
monthly_rides.index
```

```
[22]: Index(['January', 'February', 'March', 'April', 'May', 'June', 'July',  
        'August', 'September', 'October', 'November', 'December'],  
        dtype='object')
```

```
[23]: # Create a bar plot of total rides per month
```

```
plt.figure(figsize=(12,7))  
ax = sns.barplot(x=monthly_rides.index, y=monthly_rides)  
ax.set_xticklabels(month_order)  
plt.title('Ride count by month', fontsize=16);
```



Monthly rides are fairly consistent, with notable dips in the summer months of July, August, and September, and also in February.


```
[24]: # Repeat the above process, this time for rides by day
```

```
daily_rides = df['day'].value_counts()  
day_order = ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday',  
            ↪ 'Saturday', 'Sunday']  
daily_rides = daily_rides.reindex(index=day_order)  
daily_rides
```

```
[24]: Monday      2931
      Tuesday     3198
      Wednesday   3390
      Thursday    3402
      Friday      3413
      Saturday    3367
      Sunday      2998
      Name: day, dtype: int64
```

```
[25]: # Create bar plot for ride count by day

plt.figure(figsize=(12,7))
ax = sns.barplot(x=daily_rides.index, y=daily_rides)
ax.set_xticklabels(day_order)
ax.set_ylabel('Count')
plt.title('Ride count by day', fontsize=16);
```



Suprisingly, Wednesday through Saturday had the highest number of daily rides, while Sunday and Monday had the least.

```
[26]: # Repeat the process, this time for total revenue by day


day_order = ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday',
             ↪ 'Saturday', 'Sunday']
total_amount_day = df.groupby('day').sum()[['total_amount']]
total_amount_day = total_amount_day.reindex(index=day_order)
```

```
total_amount_day
```

```
[26]:          total_amount
day
Monday      49574.37
Tuesday     52527.14
Wednesday   55310.47
Thursday    57181.91
Friday      55818.74
Saturday    51195.40
Sunday      48624.06
```

```
[27]: # Create bar plot of total revenue by day

plt.figure(figsize=(12,7))
ax = sns.barplot(x=total_amount_day.index, y=total_amount_day['total_amount'])
ax.set_xticklabels(day_order)
ax.set_ylabel('Revenue (USD)')
plt.title('Total revenue by day', fontsize=16);
```



Thursday had the highest gross revenue of all days, and Sunday and Monday had the least. Interestingly, although Saturday had only 35 fewer rides than Thursday, its gross revenue was ~\$6,000 less than Thursday's—more than a 10% drop.

```
[28]: # Repeat the process, this time for total revenue by month


total_amount_month = df.groupby('month').sum()[['total_amount']]
total_amount_month = total_amount_month.reindex(index=month_order)
total_amount_month
```

```
[28]:
```

month	total_amount
January	31735.25
February	28937.89
March	33085.89
April	32012.54
May	33828.58
June	32920.52
July	26617.64
August	27759.56
September	28206.38
October	33065.83
November	30800.44
December	31261.57

```
[29]: # Create a bar plot of total revenue by month

plt.figure(figsize=(12,7))
ax = sns.barplot(x=total_amount_month.index,
                 y=total_amount_month['total_amount'])
plt.title('Total revenue by month', fontsize=16);
```



Monthly revenue generally follows the pattern of monthly rides, with noticeable dips in the summer months of July, August, and September, and also one in February.

```
[30]: # Get number of unique drop-off location IDs
```

```
df['DOLocationID'].nunique()
```

```
[30]: 216
```

```
[31]: # Calculate the mean trip distance for each drop-off location
distance_by_dropoff = df.groupby('DOLocationID').mean()[['trip_distance']]

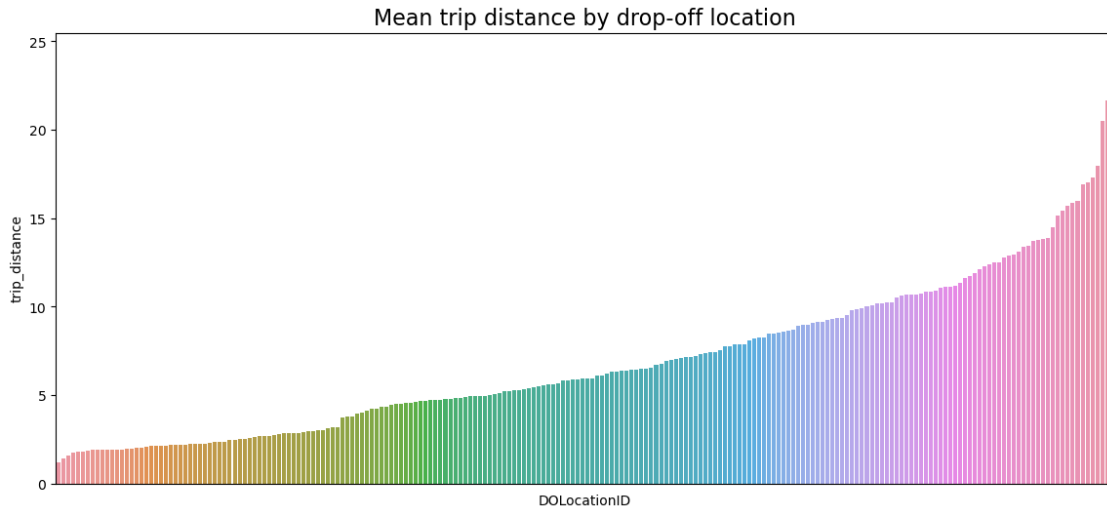
# Sort the results in descending order by mean trip distance
distance_by_dropoff = distance_by_dropoff.sort_values(by='trip_distance')
distance_by_dropoff
```

```
[31]:
```

	trip_distance
DOLocationID	
207	1.200000
193	1.390556
237	1.555494
234	1.727806
137	1.818852
...	...
51	17.310000
11	17.945000
210	20.500000
29	21.650000
23	24.275000

```
[216 rows x 1 columns]
```

```
[32]: # Create a bar plot of mean trip distances by drop-off location in ascending
      ↪order by distance
plt.figure(figsize=(14,6))
ax = sns.barplot(x=distance_by_dropoff.index,
                 y=distance_by_dropoff['trip_distance'],
                 order=distance_by_dropoff.index)
ax.set_xticklabels([])
ax.set_xticks([])
plt.title('Mean trip distance by drop-off location', fontsize=16);
```

This plot presents a characteristic curve related to the cumulative density function of a normal distribution. In other words, it indicates that the drop-off points are relatively evenly distributed over the terrain. This is good to know, because geographic coordinates were not included in this dataset, so there was no obvious way to test for the distribution of locations.

```
[33]: # 1. Generate random points on a 2D plane from a normal distribution
test = np.round(np.random.normal(10, 5, (3000, 2)), 1)
midway = int(len(test)/2) # Calculate midpoint of the array of coordinates
start = test[:midway]     # Isolate first half of array ("pick-up locations")
end = test[midway:]       # Isolate second half of array ("drop-off locations")

# 2. Calculate Euclidean distances between points in first half and second half
# of array
distances = (start - end)**2
distances = distances.sum(axis=-1)
distances = np.sqrt(distances)


# 3. Group the coordinates by "drop-off location", compute mean distance
test_df = pd.DataFrame({'start': [tuple(x) for x in start.tolist()],
                          'end': [tuple(x) for x in end.tolist()],
                          'distance': distances})
data = test_df[['end', 'distance']].groupby('end').mean()
data = data.sort_values(by='distance')

# 4. Plot the mean distance between each endpoint ("drop-off location") and all
# points it connected to
plt.figure(figsize=(14,6))
ax = sns.barplot(x=data.index,
                 y=data['distance'],
```

```

        order=data.index)
ax.set_xticklabels([])
ax.set_xticks([])
ax.set_xlabel('Endpoint')
ax.set_ylabel('Mean distance to all other points')
ax.set_title('Mean distance between points taken randomly from normal_
↳distribution');

```



The curve described by this graph is nearly identical to that of the mean distance traveled by each taxi ride to each drop-off location. This reveals that the drop-off locations in the taxi dataset are evenly distributed geographically. Note, however, that this does not mean that there was an even distribution of rides to each drop-off point. Examining this next.

```

[34]: # Check if all drop-off locations are consecutively numbered

df['DOLocationID'].max() - len(set(df['DOLocationID']))

```

[34]: 49

The above code shows that there are 49 drop-offs locations are not consecutively marked.


Hence, it is required to remove the spaces these 49 locations can introduce in histogram.

```

[35]: plt.figure(figsize=(16,4))
      # DOLocationID column is numeric, so sort in ascending order
sorted_dropoffs = df['DOLocationID'].sort_values()
      # Convert to string
sorted_dropoffs = sorted_dropoffs.astype('str')

```

```
# Plot
sns.histplot(sorted_dropoffs, bins=range(0, df['DOLocationID'].max()+1, 1))
plt.xticks([])
plt.xlabel('Drop-off locations')
plt.title('Histogram of rides by drop-off location', fontsize=16);
```



Notice that out of the 200+ drop-off locations, a disproportionate number of locations receive the majority of the traffic, while all the rest get relatively few trips. It's likely that these high-traffic locations are near popular tourist attractions like the Empire State Building or Times Square, airports, and train and bus terminals. However, it would be helpful to know the location that each ID corresponds with. Unfortunately, this is not in the data.

[]: